

Issues with current `flat_map` proposal

Document #:
Date: 2019-05-20
Project: Programming Language C++
Library Working Group
Reply-to: Denis Yaroshevskiy
<denis.yaroshevskij@gmail.com>

Contents

1	Revision history	1
2	Introduction	1
3	Listing of examples of how current design is problematic	1
4	Lack of reference implementations	2
5	Lack of zip limits the usage of <code>flat_map</code>	2
6	Performance implications of two containers on sort	2
7	How is pair of references different from <code>ref/span/string_view</code> .	3
8	Proposed actions	3
9	References	3

1 Revision history

Revision 0 Original revision of the paper for 2019 Cologne meeting.

2 Introduction

[P0429] introduces `flat_map` into the standard as a container adaptor based on two containers which requires that reference type to be a proxy object. This has been attempted on multiple occasions for other use-cases such as an infamous `vector<bool>` [VECTOR_BOOL], multi-span (decided against it) [P0009], standard audio proposal [P1386] (also decided against it following multi-span) and others. The goal of this paper is to object to inclusion of `flat_map` with same flawed design into the standard at least until the point that zip is standardized.

3 Listing of examples of how current design is problematic

In the `flat_map` paper reference type is defined as `pair<const key_type&, value_type&>`. This breaks many expectations and typical patterns for writing C++ code. (In all of the examples `decltype(it)` is `flat_map::iterator`).

```

//-----
auto x = std::move(*it); // `x` is moved from `*it`. Unless it's a `flat_map` where
                        // this creates a reference. Copy does not work either.

//-----
auto& x = *it; // Create a mutable reference to the value pointed by it. Except for flat_map where thi
              // does not compile.

//-----
[x = *it] { do_smth(x); } // Capture an element by value. Unless flat_map,
                        // in which case capture is by reference.

//-----
auto foo() { // This does not dangle unless used with flat_map's iterators.
    // ...
    return *it;
}

//-----
template <typename T>
void bar(T mine) { // bar does not modify input parameters. Unless called with flat_map referenc
    sink(std::move(mine));
}

```

4 Lack of reference implementations

There is no production (or even a complete reference) implementation for `flat_map` that is based on two containers that the author could find. All popular open source implementations (boost, folly, eastl, chromium) use a single container. The only library that we are aware of that could provide a similar experience to using two containers flat map is `zip` utility from `ranges-v3`. `zip` was not yet accepted for standardisation.

5 Lack of zip limits the usage of flat_map

- `erase_remove_if` idiom is not implementable since there is no `zip` in the standard. Also current version does not have mutating access to keys/values which also disallows this.
- One of the most typical use-cases for `flat_map` is building a buffer and then converting it into a map. Without standard `zip` we cannot use algorithms to populate such buffers.

These are very important use-cases and current proposal does not address them.

6 Performance implications of two containers on sort

The goal of separating keys and values into two arrays is to increase the lookup speed by packing keys together in the cache. However, this might have detrimental consequences for sort.

Sort is absolutely crucial for flat containers - for example most of the time spent in Chromium on every key-stroke is `flat_set` construction [[CHROMIUM_EXAMPLE](#)].

However, at least with `ranges::zip` (which is the only known example of two-container sort), it brings a significant overhead [[QUICKBENCH_SORT](#)]. At this point the author does not know whether this is due to the quality of an implementation, benchmarking artefact or a fundamental problem.

7 How is pair of references different from `ref/span/string_view`.

It might be important to clarify why such “reference like types” as `span` and `string_view` do not cause a problem while pair of references does. They actually can cause similar issues, which was the case in multi-span - if we try to return a reference to a line/column as an object that does not actually exist. The problem occurs when we break iterator/range/container concepts and not with the non-owning type itself.

8 Proposed actions

This paper suggest to postpone `flat_map` at least until the standardization of `zip`. If with `zip` the standard committee decides that proxy references are an acceptable practice in C++ than `flat_map` can use them too and the C++ community will have to learn to be extremely cautious in a much bigger number of use-cases than now. But this should be a considered decision and not something done on the back of `flat_map`. Ideally would be to find a better solution for proxy references that is intuitive and unintrusive. And then use that for `flat_map`.

9 References

- [CHROMIUM_EXAMPLE] Example where sort is important in Chromium.
https://cs.chromium.org/chromium/src/components/omnibox/browser/url_index_private_data.cc?l=657&rc1=7bfecf258f220ea375ffd376608d6ec71ca9d8ce
- [P0009] 2019. Polymorphic Multidimensional Array Reference.
<https://wg21.link/p0009>
- [P0429] Zach Laine. 2019. A Standard `flat_map`.
<https://wg21.link/p0429>
- [P1386] 2019. A Standard Audio API for C++: Motivation, Scope, and Basic Design.
<https://wg21.link/p1386>
- [QUICKBENCH_SORT] Measuring sort of two containers vs one container.
<http://quick-bench.com/V7zQvXo5R13DVMvmuRzj9GtoEMM>
- [VECTOR_BOOL] Howard Hinnant. 2012. On `vector<bool>`.
<https://howardhinnant.github.io/onvectorbool.html>